

# Designing the future of agent-server communication in RUDDER

---

Alexis Mousset - [amo@rudder.io](mailto:amo@rudder.io)

February 4, 2020

CfgMgmtCamp Ghent 2020

Context

Needs and requirements

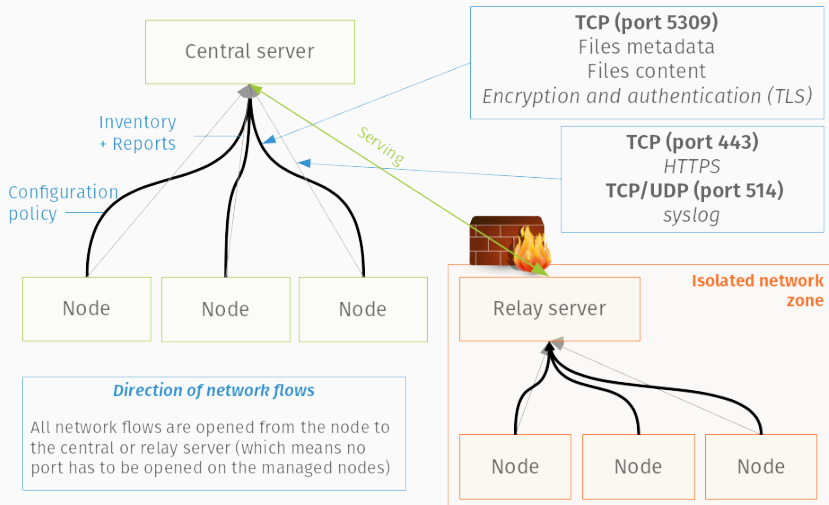
Design choices

Implementation

Perspectives

# Context

---



# Agent-server communication

- HTTP PUT for inventories
- Custom protocol in TLS for policy copy
- Standard syslog for reporting

```
R: @osquery_installation_and_configuration
@result_success
@@32377fd7-02fd-43d0-aab7-28460a91347b
@@de4435a0-b5db-401a-aba1-a685d8a937e1
@@
@@File from HTTP server
@@/etc/pki/RPM-GPG-KEY-osquery
@@2020-01-31 22:16:00+00:00
##2f4e4400-3206-4fb3-ae7f-16d1192e38ac
@#File /etc/pki/RPM-GPG-KEY-osquery is correct
```

# Reporting

R: @@Technique@@Type@@RuleId@@DirectiveId  
@@VersionId@@Component@@Key@@ExecutionTimeStamp  
##NodeId@#HumanReadableMessage

- Produced by the policies (a lot of work actually!)
- One log in syslog for each component
- Special cases for the first and last report
- The webapp only processes them when last report is received
- Parsed using `awk` for `rudder agent run output`

# Needs and requirements

---



# Reporting - Current limitations

- Reporting security
  - plain text
  - no authentication (easy to fake)
- Missing information in reports
  - Differences between expected and current state
  - Information about what has been repaired exactly
- Syslog itself
  - Requires a specific port (514)
  - Requires root access (to configure local syslog daemon)
  - Can interact with user's syslog configuration
  - Hard to debug (not much logs about syslog daemon by default)
  - Poor performance (for database insertion)

# Constraints

- Smooth transition
  - Keep compatibility with both reporting modes for several versions
  - Allow switching at any time (same data model)
- *Keep It Simple*
  - Debuggable
  - Low operation overhead
  - Use well-known technologies
- Security
  - State-of-the-art security for reporting protocol
  - Allow future homogenization
  - Focus on security for the implementation itself

# Design choices

---

# Report → Run log

- The stream of reports is useless
- Better transmitted as a single run log
- Store information by run (in a simple file)
- Easier to manage and allows lots of improvements
  - Compression (works well!)
  - Database transaction by run
- A run log is identified by: a node id, a config id and a date+time

# Improve run log

- We are missing a lot of valuable information
- "Hidden" in agent logs (`rudder agent run -i`)
- Need to `ssh` to understand anything
- We want to capture and contextualize them

```
2020-02-02T16:35:35+00:00    error: Proposed executable file '/tmp/status.sh' doesn't exist
2020-02-02T16:35:35+00:00    error: '/tmp/status.sh' promises to be executable but isn't
A| non-compliant Test_share    Command execution result  /tmp/status.sh    Execute the command /tmp/status.sh was not correct
```

# Improve run log

The problem is that we have two (isolated) information streams:

- Reports from inside of the policies
- Agent logs (errors, executed commands, various outputs, etc)

# Improve run log

- Agents are usually not designed for error management at scale, and expect human interaction.
- Nothing built-in for automated outcome analysis (what failed and what has been done)
  - no structured errors in the policies, only access to a state (**error/ok/repai**red) from inside the policy
  - no business knowledge in the logs



Use (=parse) stdout

# Improve run log

- Capture full agent output in info mode
- Parse it on the relay/server
- Associate simple logs with following contextualized log
- Works for log from the technique editor or modern techniques
- Not that good for legacy techniques: do everything then report
- Specific insertion/purging configuration for non-results logs (=simple logs)

```

2020-02-02T16:35:35+00:00    error: Proposed executable file '/tmp/status.sh' doesn't exist
2020-02-02T16:35:35+00:00    error: '/tmp/status.sh' promises to be executable but isn't
A| non-compliant Test_share    Command execution result  /tmp/status.sh    Execute the command /tmp/status.sh was not correct

```

tmp							
Execution date ▼	Status	Rule	Directive	Component	Value	Message	
⚙ Run start date: 2020-02-02 16:38:19+0000							
2020-02-02 16:38:21+0000	Noncompliant	Global configuration for all nodes	Test share	Command execution result	/tmp/status.sh	Execute the command /tmp/status.sh was not correct	
2020-02-02 16:38:21+0000	Warn	Global configuration for all nodes	Test share	Command execution result	/tmp/status.sh	'/tmp/status.sh' promises to be executable but isn't	
2020-02-02 16:38:21+0000	Warn	Global configuration for all nodes	Test share	Command execution result	/tmp/status.sh	Proposed executable file '/tmp/status.sh' doesn't exist	

# Reports authentication

- We want to authenticate reporting
- We want to stay asynchronous
- End to end validation (check signature on root server)
- We need a signature (like we do for inventories)
- Prefer a standard
- We have a hierarchical node structure

S/MIME

# HTTP

Use HTTP as it's:

- Already used for inventories (and Windows policy downloads)
- Well-known
- Easy operation and debugging (curl, etc.)
- Fast and powerful enough (even more with HTTP/2)
- Use simple file PUT (like inventories)

# Implementation

---

# Agent

- Use the existing `rudder agent run` wrapper
- Collect output, sign and compress
- Send to the server
- Retry in case of failure
- Allows back-filling compliance data



# relayd

- A new daemon that runs on all policy servers (root + relay)
- Reminder: A root server is also a relay
- Replaces relay python API
- Layer between the webapp and the nodes
- Stateless (except for history)

# Relay API

- Based on what existed since 3.2 (implemented in Python)
- Now versioned and documented
- Only listens locally
- Some endpoints behind **httd** reverse proxy
- <https://docs.rudder.io/api/relay>
- Still missing a full stats/monitoring API (prometheus?)

# Relay API

- `/system/{status, reload, info}`
- `/policies/{node-id}/rules`
- `/remote-run/nodes/{id}`
- `/shared-files/{target-id}/...`
- `/shared-folder/{path}`

# relayd

- Config files: `/opt/rudder/etc/relay`
- A new `rudder-relayd` service
- Logs to `journald`
- Part of the `rudder-server-relay` package

# Service hardening

```
User=rudder-relayd  
ProtectSystem=strict  
ReadWritePaths=/var/rudder/reports  
                /var/rudder/inventories  
                /var/rudder/shared-files  
PrivateTmp=True
```

A dedicated SELinux context

- Write access to work directories
- Read access to configuration and data files
- Connect to HTTP and postgresql ports
- Listen on port 3030
- Run the 'rudder remote run' command with sudo

# HTTP security

- Enforce TLS1.2+ everywhere (except syslog!)
- Option to check certificates in all HTTP requests
- Not directly linked
- Allows authentication of both ways
- For now, requires an existing PKI (and proper DNS setup)

# Rust

- We wanted:
  - Reliability and security
  - Maintainability (<3 strong typing)
  - Low footprint (to allow "embedded" relayd)
  - Easy packaging and deployment
- Chose Rust!



# To sum up

We added a daemon between root server and nodes to:

- Forward reports and inventories (inotify-based)
- Check, parse and store reports on root server
- Provide the file sharing API
- Provide the policy and shared-files download API (for Windows)
- Only required **httpd** and agent to synchronize data files

# Perspectives

---

# Future

- Encryption option (using S/MIME)
- Use S/MIME for inventories too
- Allow policy updates over HTTPS for full HTTP communication
- Diffs for all non-compliance or errors (including files!)
- Connected mode for reactivity and continuity
- Flexible RUDDER server distribution (container, roles, cloud, etc)
- New (virtual) agents (maybe managed from relayd)
- Check server certificate by default

Feedbacks?

Thank you!